



# Winning Space Race with Data Science

Troy Altus  
April 29, 2023

**Github Repository:**  
<https://github.com/altustd/Coursera-IBM-Applied-Data-Science-Capstone>



# Outline

---

- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

# Executive Summary

---

- Methodology and results presented herein for the IBM Data Science Professional Certificate capstone project
- Capstone project is the tenth and final course in the IBM Data Science Professional Certificate curriculum
- Project topic deals with the reusability of Space-X boosters

# Introduction

---

- The Space-X business model benefits from cost savings associated with the reuse of boosters after they are launched, recovered and serviced
  - First stage boosters have the capability to re-entering earth's atmosphere and landing by firing engines in retrograde fashion
  - Boosters are landed on terrestrial landing pads or floating barges
  - Multiple boosters can be landed from a single flight, as demonstrated by Falcon 9 heavy
- The work carried out in this project demonstrates various Data Science techniques based on actual flight data

---

# Methodology

# Methodology Overview

---

- Data collection
- Data wrangling
- Exploratory data analysis (EDA) using visualization and SQL
- Interactive visual analytics using Folium and Plotly Dash
- Predictive analysis using classification models

# Methodology: Data Collection

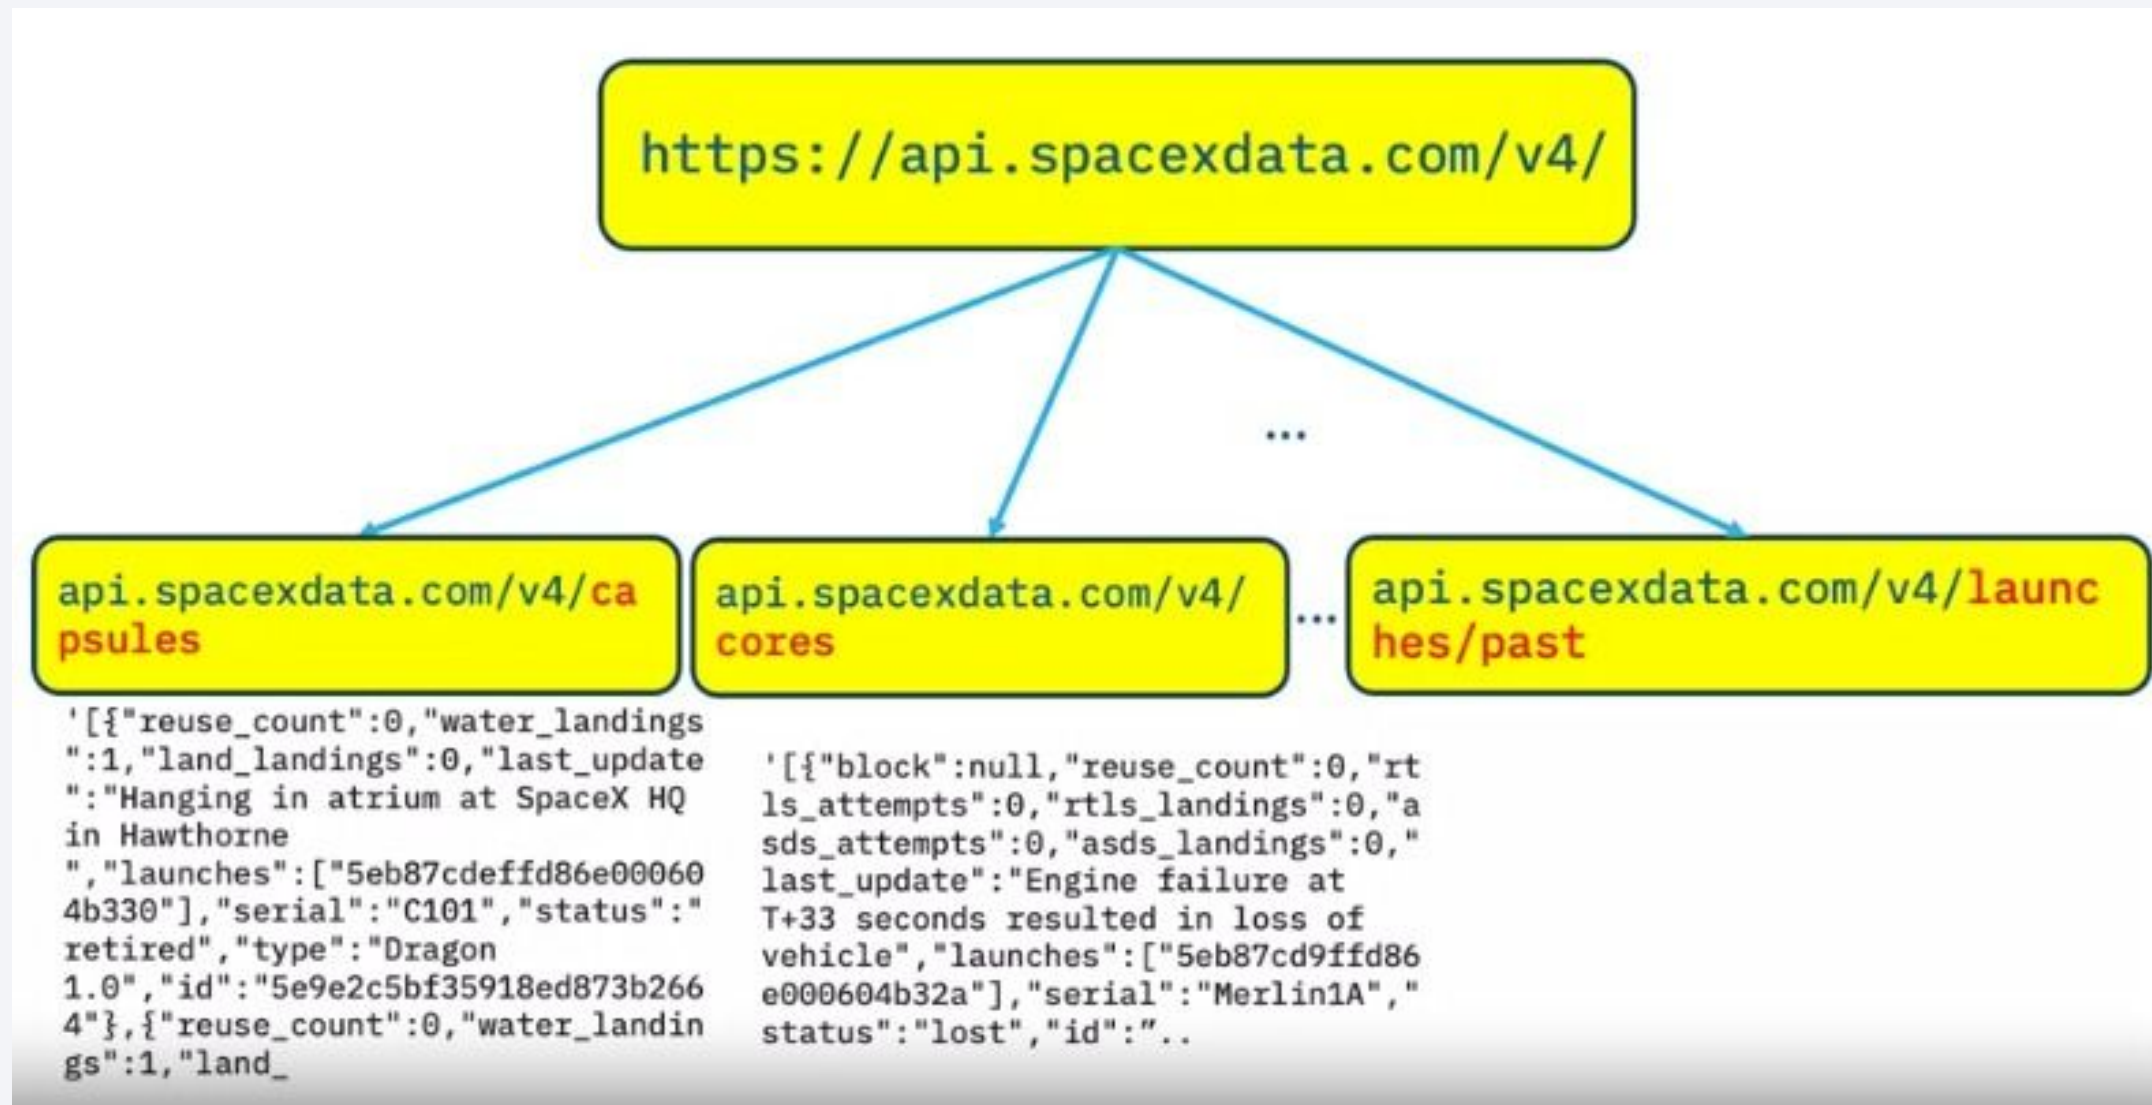
---

- Goal is to transform raw data into a clean dataset by wrangling data using an API, filtering and sampling the data, and dealing with nulls
- SpaceX REST API gives data about launches
  - Type of rocket used
  - Payload delivered
  - Launch specifications
  - Landing specification
  - Landing outcome
- Additional Falcon 9 launch data obtained via web scraping related Wiki pages
  - Python BeautifulSoup package used to web scrape HTML tables containing Falcon 9 launch records
  - Data is parsed and converted to a Pandas data frame for further visualization and analysis
- Original data set (Falcon 1 + Falcon 9) is filtered to only include Falcon 9 launches
- Null values are addressed using realistic assumptions
  - PayloadMass null values replaced by mean value
  - LandingPad null values are dealt with via one hot encoding



# Methodology: Data Collection – SpaceX API

## Top-level API Flowchart





# Methodology: Data Collection – SpaceX API

## API endpoint call

- DefinesSpecific url pointing to historical launch data
- Perform “get” request from the requests library
- Response is in the form of a list of .json objects

```
[6]: spacex_url="https://api.spacexdata.com/v4/launches/past"
```

```
[7]: response = requests.get(spacex_url)
```

Check the content of the response

```
[29]: # Uncomment for result.
```

```
print(response.content)
```

```
b' [{"fairings": {"reused": false, "recovery_attempt": false, "recovered": false, "ships": []}, "links": {"patch": {"small": "https://images2.imgbox.com/94/f2/NN6Ph45r_o.png", "large": "https://images2.imgbox.com/5b/02/QcxHUB5V_o.png"}, "reddit": {"campaign": null, "launch": null, "media": null, "recovery": null}, "flickr": {"small": [], "original": []}, "presskit": null, "webcast": "https://www.youtube.com/watch?v=0a_00nJ_Y88", "youtube_id": "0a_00nJ_Y88", "article": "https://www.space.com/2196-spacex-inaugural-falcon-1-rocket-lost-launch.html", "wikipedia": "https://en.wikipedia.org/wiki/DemoSat"}, "static_fire_date_utc": "2006-03-17T00:00:00.000Z", "static_fire_date_unix": 1142553600, "net": false, "window": 0, "rocket": "5e9d0d95eda69955f709d1eb", "success": false, "failures": [{"time": 33, "altitude": null, "reason": "merlin engine failure"}], "details": "Engine failure at 33 seconds and loss of vehicle", "crew": [], "ships": [], "capsules": [], "payloads": ["5eb0e4b5b6c3bb0006eeb1e1"], "launchpad": "5e9e4502f5090995de566f86", "flight_number": 1, "name": "FalconSat", "date_utc": "2006-03-24T22:30:00.000Z", "date_unix": 1143239400, "date_local": "2006-03-25T10:30:00+12:00", "date_precision": "hour", "upcoming": false, "cores": [{"core": "5e9e289df35918033d3b2623", "flight": 1, "gridfins": false, "legs": false, "reused": false, "landing_attempt": false, "landing_success": null, "landing_type": null, "landpad": null}], "auto_update": true, "tbd": false, "launch_library_id": null, "id": "5eb87cd9ffd86e000604b32a"}, {"fairings": {"reused": false, "recovery_attempt": false, "recovered": false, "ships": []}, "links": {"patch": {"small": "https://images2.imgbox.com/f9/4a/ZboXReNb_o.png", "large": "https://images2.imgbox.com/80/a2/bkWotCIS_o.png"}, "reddit": {"campaign": null, "launch": null, "media": null, "recovery": null}, "flickr": {"small": [], "original": []}, "presskit": null, "webcast": "https://www.youtube.com/watch?v=Lk4zQ2wP-Nc", "youtube_id": "Lk4zQ2wP-Nc", "article": "https://www.space.com/3590-spacex-falcon-1-rocket-fails-reach-orbit.html", "wikipedia": "https://en.wikipedia.org/wiki/DemoSat"}}, {"static_fire_date_utc": null, "static_fire_date_unix": null, "net": false, "window": 0, "rocket": "5e9d0d95eda69955f709d1eb", "success": false, "failures": [{"time": 33, "altitude": null, "reason": "merlin engine failure"}], "details": "Engine failure at 33 seconds and loss of vehicle", "crew": [], "ships": [], "capsules": [], "payloads": ["5eb0e4b5b6c3bb0006eeb1e1"], "launchpad": "5e9e4502f5090995de566f86", "flight_number": 1, "name": "FalconSat", "date_utc": "2006-03-24T22:30:00.000Z", "date_unix": 1143239400, "date_local": "2006-03-25T10:30:00+12:00", "date_precision": "hour", "upcoming": false, "cores": [{"core": "5e9e289df35918033d3b2623", "flight": 1, "gridfins": false, "legs": false, "reused": false, "landing_attempt": false, "landing_success": null, "landing_type": null, "landpad": null}], "auto_update": true, "tbd": false, "launch_library_id": null, "id": "5eb87cd9ffd86e000604b32a"}]
```

# Methodology: Data Collection – SpaceX API

## API Data Wrangling

- Individual .json objects from API call are converted to a DataFrame using “json\_normalize” function
- This allows us to convert to a flat table

```
[11]: # Use json_normalize method to convert the json result into a dataframe
respjson = response.json()
data = pd.json_normalize(respjson)
```

Using the dataframe `data` print the first 5 rows

```
[12]: # Get the head of the dataframe
data.head(5)
```

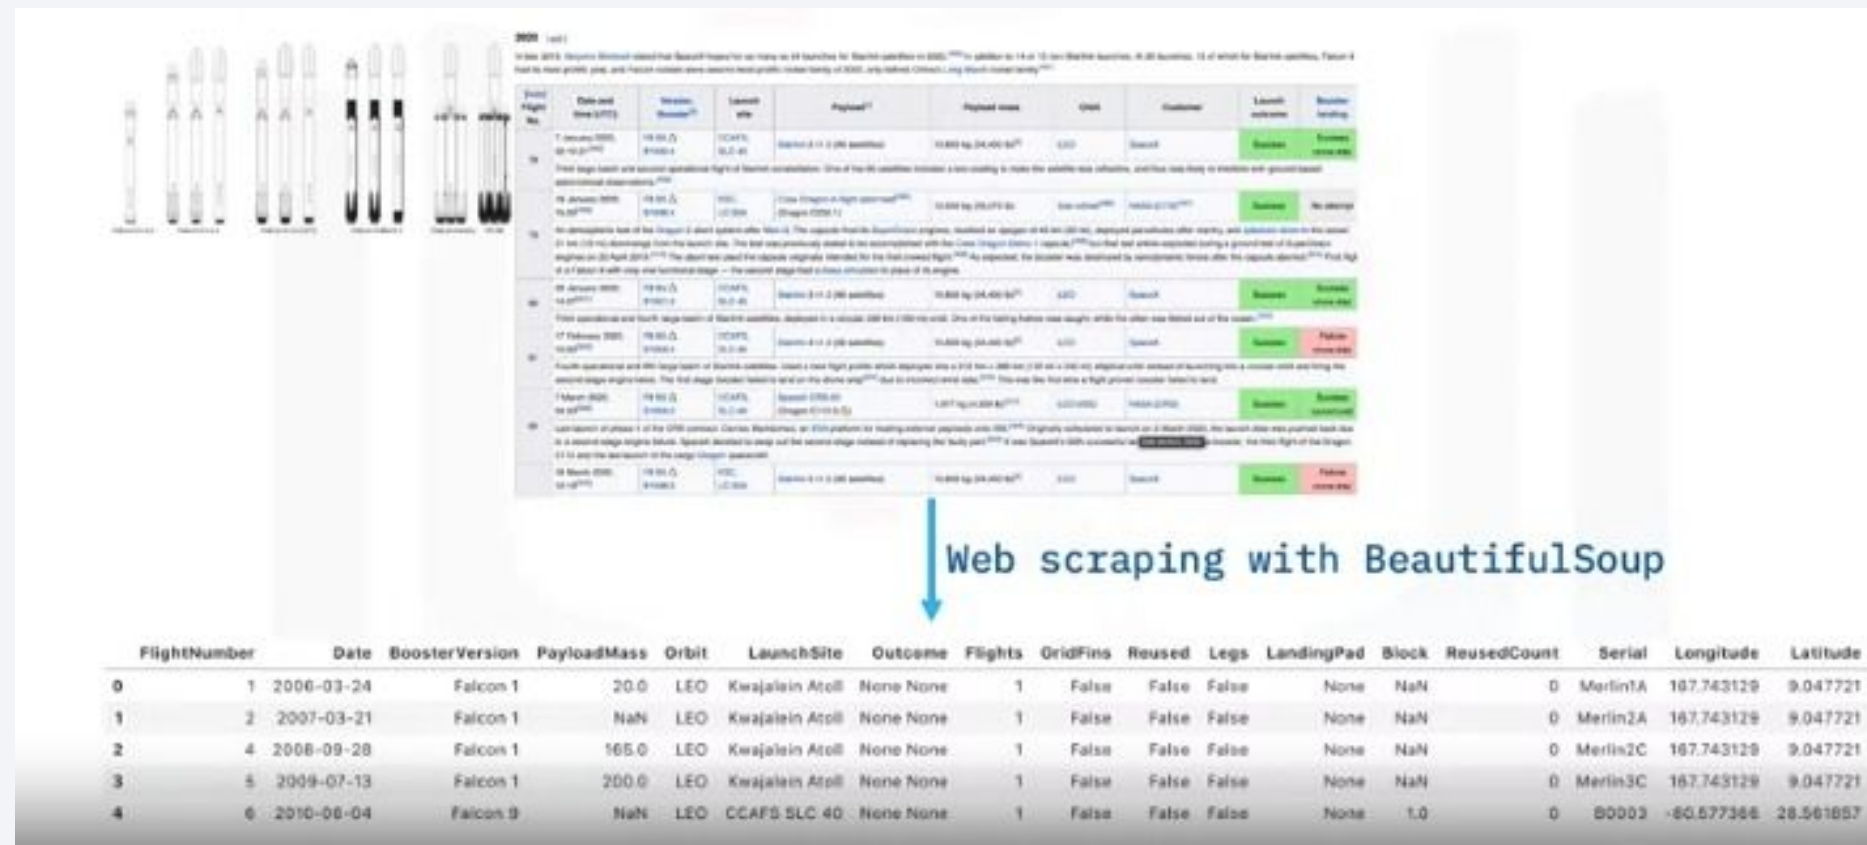
```
[12]:
```

	static_fire_date_utc	static_fire_date_unix	net	window	rocket	success	failures	details	crew	ships	capsules	pay
0	2006-03-17T00:00:00.000Z	1.142554e+09	False	0.0	5e9d0d95eda69955f709d1eb	False	[[{'time': 33, 'altitude': None, 'reason': 'merlin engine failure'}]]	Engine failure at 33 seconds and loss of vehicle	[]	[]	[]	[5eb0e4b5b6c3bb0006ee
1	None	NaN	False	0.0	5e9d0d95eda69955f709d1eb	False	[[{'time': 301, 'altitude': 289, 'reason': 'harmonic oscillation leading to premature engine shutdown at T+7 min 30 s. Failed	Successful first stage burn and transition to second stage, maximum altitude 289 km, Premature engine shutdown	[]	[]	[]	[5eb0e4b6b6c3bb0006ee

# Methodology: Data Collection - Web Scrapping

## Web Scrapping Related Wiki Pages

- Python BeautifulSoup library
- Tabular data is parsed and converted to a Pandas DataFrame



The image shows a screenshot of a Wikipedia page for the Falcon 1 rocket. A table of flight data is visible, with columns for Flight Number, Date, Booster Version, Payload Mass, Orbit, Launch Site, Outcome, Flights, Grid Fins, Reused, Legs, Landing Pad, Block, Reused Count, Serial, Longitude, and Latitude. An arrow points from the table to a Pandas DataFrame below.

Web scrapping with BeautifulSoup

FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block	ReusedCount	Serial	Longitude	Latitude
0	1 2006-03-24	Falcon 1	20.0	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin1A	167.743129	9.047721
1	2 2007-03-21	Falcon 1	NaN	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin2A	167.743129	9.047721
2	4 2008-09-28	Falcon 1	165.0	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin2C	167.743129	9.047721
3	5 2009-07-13	Falcon 1	200.0	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin3C	167.743129	9.047721
4	6 2010-06-04	Falcon 9	NaN	LEO	CCAFS SLC 40	None None	1	False	False	False	None	1.0	0	B0003	-80.677366	28.561857

# Methodology: Data Wrangling

---

API is used a second time to target a different endpoint

- Call to the first API endpoint provides ID number, not actual data
- Pre-defined functions provided to create the data set
- Data is stored in lists

Function	Targets	Endpoint
<code>getBoosterVersion</code>	→	Rockets URL: <a href="https://api.spacexdata.com/v4/rock">https://api.spacexdata.com/v4/rock</a>
<code>getLaunchSite</code>	→	Launchpads URL: <a href="https://api.spacexdata.com/v4/laur">https://api.spacexdata.com/v4/laur</a>
<code>getPayloadData</code>	→	Payloads URL: <a href="https://api.spacexdata.com/v4/payl">https://api.spacexdata.com/v4/payl</a>
<code>getCoreData</code>	→	<code>getCoreData</code> URL: <a href="https://api.spacexdata.com/v4/core">https://api.spacexdata.com/v4/core</a>



# Methodology: Sampling/Filtering

Data includes Falcon 1 booster, but we are only concerned with Falcon 9 for this study

- Filter the data dataframe using the BoosterVersion column to only keep the Falcon 9 launches.
- Save the filtered data to a new dataframe called data\_falcon9.

	FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block	ReusedCount	Serial	Longitude	Latitude
0	1	2006-03-24	Falcon 1	20.0	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin1A	167.743129	9.047721
1	2	2007-03-21	Falcon 1	NaN	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin2A	167.743129	9.047721
2	4	2008-09-28	Falcon 1	165.0	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin2C	167.743129	9.047721
3	5	2009-07-13	Falcon 1	200.0	LEO	Kwajalein Atoll	None None	1	False	False	False	None	NaN	0	Merlin3C	167.743129	9.047721
4	6	2010-06-04	Falcon 9	NaN	LEO	CCAFS SLC 40	None None	1	False	False	False	None	1.0	0	B0003	-80.577366	28.561857



```
[25]: data_falcon9.head()
```

	FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block	ReusedCount	Serial	Longitude	Latitude
4	6	2010-06-04	Falcon 9	NaN	LEO	CCSFS SLC 40	None None	1	False	False	False	None	1.0	0	B0003	-80.5	
5	8	2012-05-22	Falcon 9	525.0	LEO	CCSFS SLC 40	None None	1	False	False	False	None	1.0	0	B0005	-80.5	
6	10	2013-03-01	Falcon 9	677.0	ISS	CCSFS SLC 40	None None	1	False	False	False	None	1.0	0	B0007	-80.5	
7	11	2013-09-29	Falcon 9	500.0	PO	VAFB SLC 4E	False Ocean	1	False	False	False	None	1.0	0	B1003	-120.6	
8	12	2013-12-03	Falcon 9	3170.0	GTO	CCSFS SLC 40	None None	1	False	False	False	None	1.0	0	B1004	-80.5	


# Methodology: Sampling/Filtering

## Addressing Null Values

- Replace null values for payload mass with mean values
- Leave the column “Landing Pad” with null values as it is represented when a landing pad is not used. This will be dealt with when using One Hot encoding later on.

```
data_falcon9.isnull().sum()
```

FlightNumber	0
Date	0
BoosterVersion	0
PayloadMass	5
Orbit	0
LaunchSite	0
Outcome	0
Flights	0
GridFins	0
Reused	0
Legs	0
LandingPad	26
Block	0
ReusedCount	0
Serial	0
Longitude	0
Latitude	0
dtype:	int64



```
data_falcon9.isnull().sum()
```

FlightNumber	0
Date	0
BoosterVersion	0
PayloadMass	0
Orbit	0
LaunchSite	0
Outcome	0
Flights	0
GridFins	0
Reused	0
Legs	0
LandingPad	26
Block	0
ReusedCount	0
Serial	0
Longitude	0
Latitude	0
dtype:	int64



# Methodology: Exploratory Data Analysis

---

Exploratory data analysis (EDA) is used by data scientists to analyze and investigate data sets and summarize their main characteristics, often employing data visualization methods. It helps determine how best to manipulate data sources to get the answers you need, making it easier for data scientists to discover patterns, spot anomalies, test a hypothesis, or check assumptions.

EDA is primarily used to see what data can reveal beyond the formal modeling or hypothesis testing task and provides a better understanding of data set variables and the relationships between them. It can also help determine if the statistical techniques you are considering for data analysis are appropriate. Originally developed by American mathematician John Tukey in the 1970s, EDA techniques continue to be a widely used method in the data discovery process today.

REF: <https://www.ibm.com/topics/exploratory-data-analysis>

# Methodology: Exploratory Data Analysis

---

For this Capstone project, Exploratory Data Analysis was accomplished through completion of two labs:

1. Exploratory Data Analysis using SQL
  - a. IBM Skills Network Labs version used
  - b. Download datasets and store in table
  - c. Connect to the database
  - d. Write and execute SQL queries to solve the assignment tasks.
  
2. Exploratory Data Analysis with visualization and Feature Engineering using Pandas and Matplotlib
  - a. Visualize the relationship between different parameters
  - b. Visualize the launch success yearly trend
  - c. Create dummy variables to categorical columns

# Methodology: EDA with SQL

---

A dataset of historical launches was used to determine first stage landing success. Data included are as follows:

Date	Payload mass
Time (UTC)	Orbit
Booster version	Customer
Launch site	Mission outcome
Payload	Landing outcome

The general Exploratory Data Analysis steps were:

- Understand the SpaceX DataSet
- Load the dataset into the corresponding table in a Db2 database
- Execute SQL queries to answer assignment questions

# Methodology: EDA with SQL

## SQL Queries

---

Per lab instructions, the following SQL queries were carried out:

Task 1. Display the names of the unique launch sites in the space mission

Task 2. Display 5 records where launch sites begin with the string 'CCA'

Task 3. Display the total payload mass carried by boosters launched by NASA (CRS)

Task 4. Display average payload mass carried by booster version F9 v1.1

Task 5. List the date when the first succesful landing outcome in ground pad was acheived

Task 6. List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000

Task 7. List the total number of successful and failure mission outcomes

Task 8. List the names of the booster\_versions which have carried the maximum payload mass. Use a subquery

Task 9. List the records which will display the month names, failure landing\_outcomes in drone ship ,booster versions, launch\_site for the months in year 2015

**Jupyter Notebook can be found here:**

# Methodology: EDA with Visualization and Feature Engineering

---

Historical SpaceX data was used to identify which launch parameters were associated with successful first stage landings. Data for these parameters were isolated and prepared for use in a predictive model (future step in Capstone project)

Pandas and Seaborn were used to graphically explore the SpaceX launch data.

- Seaborn is a data visualization tool built on Matplotlib.
- One-hot encoding was used to convert categorical variables into dummy numerical variables for easy use in a machine learning algorithm.

Observations were as follows: The more massive the payload, the less likely the first stage will return

- Success rate increased as more flights were attempted
- No launches of payload mass exceeding 10,000 kg were attempted at Vandenberg Air Force Base (VAFB-SLC).
- The LEO success rate appears related to the number of flights; on the other hand, there seems to be no relationship between flight number when in GTO orbit.
- With heavy payloads, positive landing rates are higher for polar, LEO and ISS. This was not distinguishable for GTO launches.

## Methodology: EDA with Visualization and Feature Engineering

---

Per lab instructions, the Pandas and Matplotlib were used to carry out the following tasks:

Task 1: Visualize the relationship between Flight Number and Launch Site

Task 2: Visualize the relationship between Payload and Launch Site

Task 3: Visualize the relationship between success rate of each orbit type

Task 4: Visualize the relationship between FlightNumber and Orbit type

Task 5: Visualize the relationship between Payload and Orbit type

Task 6: Visualize the launch success yearly trend

Task 7: Create dummy variables to categorical columns

Task 8: Cast all numeric columns to float64

**Jupyter Notebook can be found here:**



# Methodology: Interactive Maps with Folium

---

Folium was used for interactive plotting:

- Plot of east coast and west coast US launch sites and Johnson Space Center in Houston
- Launch sites with launch results (successes and failures)
- Launch site distance to points of interest: major roads, railways, coastline and closest city

From the plots, the following observations can be made:

1. Launch sites are close in proximity to essential infrastructure such as railways and major surface roads
2. Launch sites are reasonably far away from populated areas
3. Launch sites are near the coastline so flight path is over ocean

**Jupyter Notebook can be found here:**

# Methodology: Dashboarding with Plotly Dash

---

## Plotly – An Overview

---

- Interactive, open-source plotting library
- Supports over 40 unique chart types
- Includes chart types like statistical, financial, maps, scientific, and 3-dimensional
- Visualizations can be displayed in Jupyter notebook, saved to HTML files, or can be used in developing Python-built web applications

# Methodology: Dashboarding with Plotly Dash

---

## Plotly Sub-modules

---

- Plotly Graph Objects: Low-level interface to figures, traces, and layout

`plotly.graph_objects.Figure`

- Plotly Express: High-level wrapper

# Methodology: Dashboarding with Plotly Dash

---

## **Dash – An Overview**

---


- Open-source User Interface Python library from Plotly
- Easy to build GUI
- Declarative and Reactive
- Rendered in web browser and can be deployed to servers
- Inherently cross-platform and mobile ready

# Methodology: Predictive Analysis

---

## Build a Machine Learning Pipeline

- Predict whether first stage of Falcon 9 will land successfully



```
|from sklearn.model_selection import GridSearchCV
```

---

# Results



# Results: Exploratory Data Analysis

## SQL Queries

### Task 1

Display the names of the unique launch sites in the space mission

```
[9]: %sql SELECT DISTINCT(launch_site) FROM spacextbl;
```

```
* sqlite:///my_data1.db
```

Done.

```
[9]: launch_site
```

```
CCAFS LC-40
```

```
VAFB SLC-4E
```

```
KSC LC-39A
```

```
CCAFS SLC-40
```

### Task 2

Display 5 records where launch sites begin with the string 'CCA'

```
[23]: %sql SELECT * FROM spacextbl WHERE launch_site LIKE 'CCA%' LIMIT 5;
```

```
* sqlite:///my_data1.db
```

Done.

```
[23]:
```

	id	date	time_utc	booster_version	launch_site	payload	payload_mass_kg	orbit	customer	mission_outcome	landing_outcome
0	04-06-2010	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit		0	LEO	SpaceX	Success	Failure (parachute)
1	08-12-2010	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of Brouere cheese		0	LEO (ISS)	NASA (COTS) NRO	Success	Failure (parachute)
2	22-05-2012	07:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2		525	LEO (ISS)	NASA (COTS)	Success	No attempt
3	08-10-2012	00:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1		500	LEO (ISS)	NASA (CRS)	Success	No attempt
4	01-03-2013	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2		677	LEO (ISS)	NASA (CRS)	Success	No attempt

# Results: Exploratory Data Analysis

## SQL Queries

---

### Task 3

Display the total payload mass carried by boosters launched by NASA (CRS)

```
[24]: %sql SELECT SUM(payload_mass_kg) AS total_pl_mass_all_NASACRS FROM spacextbl WHERE customer = 'NASA (CRS)';
```

```
* sqlite:///my_data1.db  
Done.
```

```
[24]: total_pl_mass_all_NASACRS  
-----  
45596
```

### Task 4

Display average payload mass carried by booster version F9 v1.1

```
[26]: %sql SELECT AVG(payload_mass_kg) FROM spacextbl WHERE booster_version LIKE 'F9 v1.1%'
```

```
* sqlite:///my_data1.db  
Done.
```

```
[26]: AVG(payload_mass_kg)  
-----  
2534.6666666666665
```

# Results: Exploratory Data Analysis

## SQL Queries

### Task 5

List the date when the first succesful landing outcome in ground pad was acheived.

*Hint: Use min function*

```
13]: %sql SELECT MIN(id), date, landing_outcome FROM spacextbl WHERE landing_outcome = "Success (ground pad)"
```



```
* sqlite:///my_data1.db
```

Done.

```
13]:
```

MIN(id)	date	landing_outcome
19	22-12-2015	Success (ground pad)

### Task 6

List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000

```
[14]: %%sql SELECT booster_version FROM spacextbl  
WHERE landing_outcome = 'Success (drone ship)' AND payload_mass_kg > 4000 AND payload_mass_kg < 6000;
```

```
* sqlite:///my_data1.db
```

Done.

```
[14]:
```

booster_version
F9 FT B1022
F9 FT B1026
F9 FT B1021.2
F9 FT B1031.2

# Results: Exploratory Data Analysis

## SQL Queries

### Task 7

List the total number of successful and failure mission outcomes

```
[15]: %%sql
SELECT DISTINCT(mission_outcome), COUNT(mission_outcome) AS counts
FROM spacextbl
GROUP BY mission_outcome;
```

```
* sqlite:///my_data1.db
Done.
```

```
[15]:
```

mission_outcome	counts
Failure (in flight)	1
Success	98
Success	1
Success (payload status unclear)	1

Clean up results using CASE

```
[16]: %%sql
SELECT
    SUM(CASE WHEN mission_outcome LIKE 'Success%' THEN 1 else 0 END) AS overall_mission_success,
    SUM(CASE WHEN mission_outcome LIKE 'Failure%' THEN 1 else 0 END) AS overall_mission_failure,
    COUNT(*) AS total
FROM spacextbl;
```

```
* sqlite:///my_data1.db
Done.
```

```
[16]:
```

overall_mission_success	overall_mission_failure	total
100	1	101

# Results: Exploratory Data Analysis

## SQL Queries

### ▼ Task 8

List the names of the booster\_versions which have carried the maximum payload mass. Use a subquery [🔗](#)

```
[17]: %%sql SELECT DISTINCT(booster_version), payload_mass_kg
      FROM spacextbl
      WHERE payload_mass_kg = (SELECT MAX(payload_mass_kg)
      FROM spacextbl)
```

```
* sqlite:///my_data1.db
Done.
```

```
[17]: booster_version  payload_mass_kg
```

booster_version	payload_mass_kg
F9 B5 B1048.4	15600
F9 B5 B1049.4	15600
F9 B5 B1051.3	15600
F9 B5 B1056.4	15600
F9 B5 B1048.5	15600
F9 B5 B1051.4	15600
F9 B5 B1049.5	15600
F9 B5 B1060.2	15600
F9 B5 B1058.3	15600
F9 B5 B1051.6	15600
F9 B5 B1060.3	15600
F9 B5 B1049.7	15600

# Results: Exploratory Data Analysis

## SQL Queries

---

### ▼ Task 9

List the records which will display the month names, failure landing\_outcomes in drone ship ,booster versions, launch\_site for the months in year 2015.

**Note: SQLite does not support monthnames. So you need to use substr(Date, 4, 2) as month to get the months and substr(Date,7,4)='2015' for year.**

```
[18]: %%sql SELECT COUNT(mission_outcome) AS failed_mission
      FROM spacextbl
      WHERE mission_outcome LIKE 'Failure%';
```

```
* sqlite:///my_data1.db
Done.
```

```
[18]: failed_mission
```

```
1
```

```
[19]: %%sql SELECT id, date, landing_outcome, booster_version, launch_site
      FROM spacextbl
      WHERE landing_outcome = "Failure (drone ship)" AND date LIKE "%2015"
      ORDER BY id
```

```
* sqlite:///my_data1.db
Done.
```

```
[19]:
```

	id	date	landing_outcome	booster_version	launch_site
	13	10-01-2015	Failure (drone ship)	F9 v1.1 B1012	CCAFS LC-40
	16	14-04-2015	Failure (drone ship)	F9 v1.1 B1015	CCAFS LC-40



# Results: Exploratory Data Analysis

## SQL Queries

### Task 10

Rank the count of successful landing\_outcomes between the date 04-06-2010 and 20-03-2017 in descending order.

```
[20]: %%sql SELECT id, date, landing_outcome
      FROM spacextbl
      WHERE landing_outcome LIKE "Success%"
      ORDER BY id DESC
```

\* sqlite:///my\_data1.db

Done.

```
[20]:
```

	id	date	landing_outcome
	100	06-12-2020	Success
	99	25-11-2020	Success
	98	21-11-2020	Success
	97	16-11-2020	Success
	96	05-11-2020	Success
	95	24-10-2020	Success
	94	18-10-2020	Success
	93	06-10-2020	Success
	92	03-09-2020	Success
	91	30-08-2020	Success
	90	18-08-2020	Success

89	07-08-2020	Success
88	20-07-2020	Success
87	30-06-2020	Success
86	13-06-2020	Success
85	04-06-2020	Success
84	30-05-2020	Success
83	22-04-2020	Success
81	07-03-2020	Success
79	29-01-2020	Success
77	07-01-2020	Success
76	17-12-2019	Success
75	05-12-2019	Success
74	11-11-2019	Success
72	25-07-2019	Success
71	12-06-2019	Success
70	24-05-2019	Success

69	04-05-2019	Success
68	02-03-2019	Success
67	22-02-2019	Success
66	11-01-2019	Success
63	03-12-2018	Success
62	15-11-2018	Success
61	08-10-2018	Success
60	10-09-2018	Success
59	07-08-2018	Success
58	25-07-2018	Success
57	22-07-2018	Success
53	11-05-2018	Success (drone ship)
52	18-04-2018	Success (drone ship)
46	08-01-2018	Success (ground pad)
44	15-12-2017	Success (ground pad)
43	30-10-2017	Success (drone ship)

42	11-10-2017	Success (drone ship)
41	09-10-2017	Success (drone ship)
40	07-09-2017	Success (ground pad)
39	24-08-2017	Success (drone ship)
38	14-08-2017	Success (ground pad)
36	25-06-2017	Success (drone ship)
35	23-06-2017	Success (drone ship)
34	03-06-2017	Success (ground pad)
32	01-05-2017	Success (ground pad)
31	30-03-2017	Success (drone ship)
29	19-02-2017	Success (ground pad)
28	14-01-2017	Success (drone ship)
27	14-08-2016	Success (drone ship)
26	18-07-2016	Success (ground pad)
24	27-05-2016	Success (drone ship)
23	06-05-2016	Success (drone ship)
22	08-04-2016	Success (drone ship)
19	22-12-2015	Success (ground pad)

# Results: Exploratory Data Analysis

## SQL Queries

```
[21]: %%sql SELECT date, payload, landing_outcome
      FROM spacextbl
      WHERE landing_outcome LIKE 'Success%' AND date > '04-06-2010' AND date < '20-03-2017';
```

```
* sqlite:///my_data1.db
```

Done.

```
[21]:
```

date	payload	landing_outcome
08-04-2016	SpaceX CRS-8	Success (drone ship)
06-05-2016	JCSAT-14	Success (drone ship)
18-07-2016	SpaceX CRS-9	Success (ground pad)
14-08-2016	JCSAT-16	Success (drone ship)
14-01-2017	Iridium NEXT 1	Success (drone ship)
19-02-2017	SpaceX CRS-10	Success (ground pad)
14-08-2017	SpaceX CRS-12	Success (ground pad)
07-09-2017	Boeing X-37B OTV-5	Success (ground pad)
09-10-2017	Iridium NEXT 3	Success (drone ship)
11-10-2017	SES-11 / EchoStar 105	Success (drone ship)
15-12-2017	SpaceX CRS-13	Success (ground pad)
08-01-2018	Zuma	Success (ground pad)
18-04-2018	Transiting Exoplanet Survey Satellite (TESS)	Success (drone ship)
11-05-2018	Bangabandhu-1	Success (drone ship)
07-08-2018	Merah Putih	Success
10-09-2018	Telstar 18V / Apstar-5C	Success

08-10-2018	SAOCOM 1A	Success
15-11-2018	Es hail 2	Success
11-01-2019	Iridium NEXT-8	Success
12-06-2019	RADARSAT Constellation, SpaceX CRS-18	Success
11-11-2019	Starlink 1 v1.0, SpaceX CRS-19	Success
05-12-2019	SpaceX CRS-19, JCSat-18 / Kacific 1	Success
17-12-2019	JCSat-18 / Kacific 1, Starlink 2 v1.0	Success
07-01-2020	Starlink 2 v1.0, Crew Dragon in-flight abort test	Success
07-03-2020	SpaceX CRS-20, Starlink 5 v1.0	Success
04-06-2020	Starlink 7 v1.0, Starlink 8 v1.0	Success
13-06-2020	Starlink 8 v1.0, SkySats-16, -17, -18, GPS III-03	Success
07-08-2020	Starlink 9 v1.0, SXRS-1, Starlink 10 v1.0	Success
18-08-2020	Starlink 10 v1.0, SkySat-19, -20, -21, SAOCOM 1B	Success
06-10-2020	Starlink 12 v1.0, Starlink 13 v1.0	Success
18-10-2020	Starlink 13 v1.0, Starlink 14 v1.0	Success
05-11-2020	GPS III-04 , Crew-1	Success
16-11-2020	Crew-1, Sentinel-6 Michael Freilich	Success
06-12-2020	SpaceX CRS-21	Success

# Results: Exploratory Data Analysis

## SQL Queries

---

```
22]: %%sql
SELECT landing_outcome, COUNT(landing_outcome) AS Total_Landing_Outcome
FROM spacextbl
WHERE landing_outcome = 'Failure (drone ship)' OR landing_outcome = 'Success (ground pad)'
AND date BETWEEN '2010-06-04' AND '2017-03-20'
GROUP BY landing_outcome
ORDER BY Total_Landing_Outcome DESC;
```

```
* sqlite:///my_data1.db
Done.
```

```
22]:
```

landing_outcome	Total_Landing_Outcome
Failure (drone ship)	5

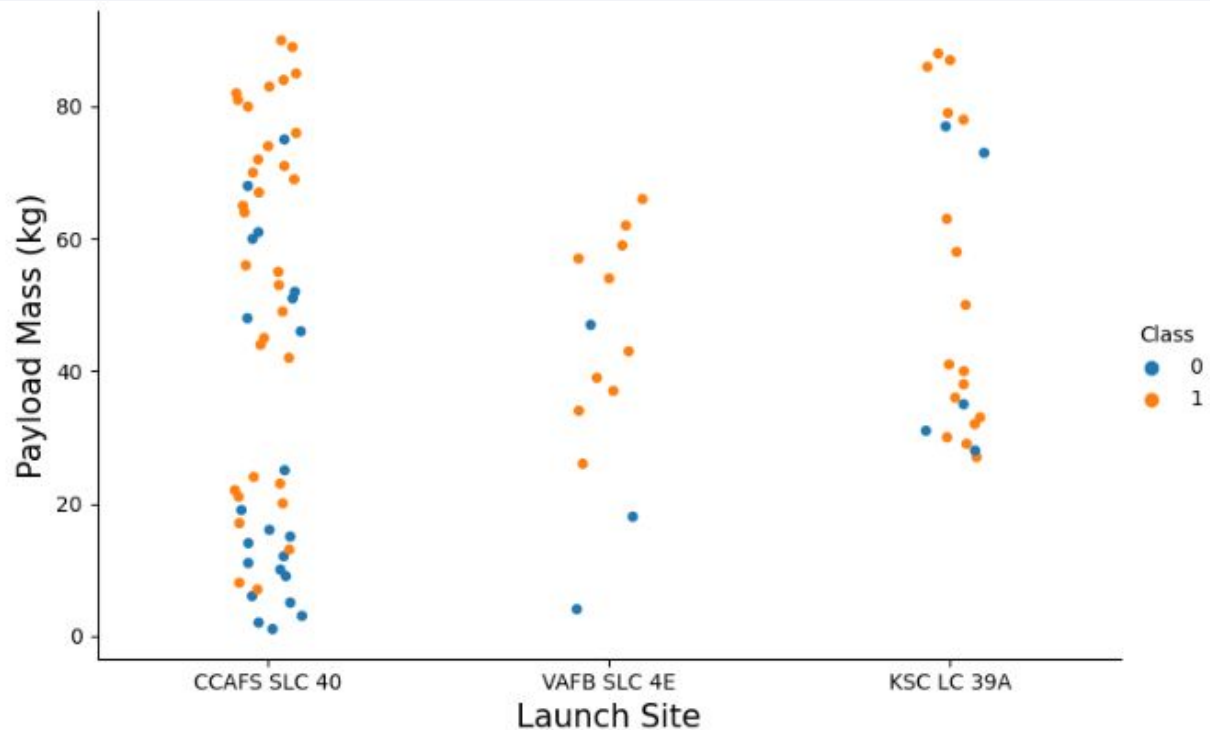
# Results: Exploratory Data Analysis

## Visualization and Feature Engineering

### TASK 1: Visualize the relationship between Flight Number and Launch Site

Use the function `catplot` to plot `FlightNumber` vs `LaunchSite`, set the parameter `x` parameter to `FlightNumber`, set the `y` to `Launch Site` and set the parameter `hue` to `'class'`

```
[20]: # Plot a scatter point chart with x axis to be Flight Number and y axis to be the launch site, and hue to be the class value
sns.catplot(y="FlightNumber", x="LaunchSite", hue='Class', data=df, aspect = 1.5)
plt.xlabel("Launch Site", fontsize=15)
plt.ylabel("Payload Mass (kg)", fontsize=15)
plt.show()
```



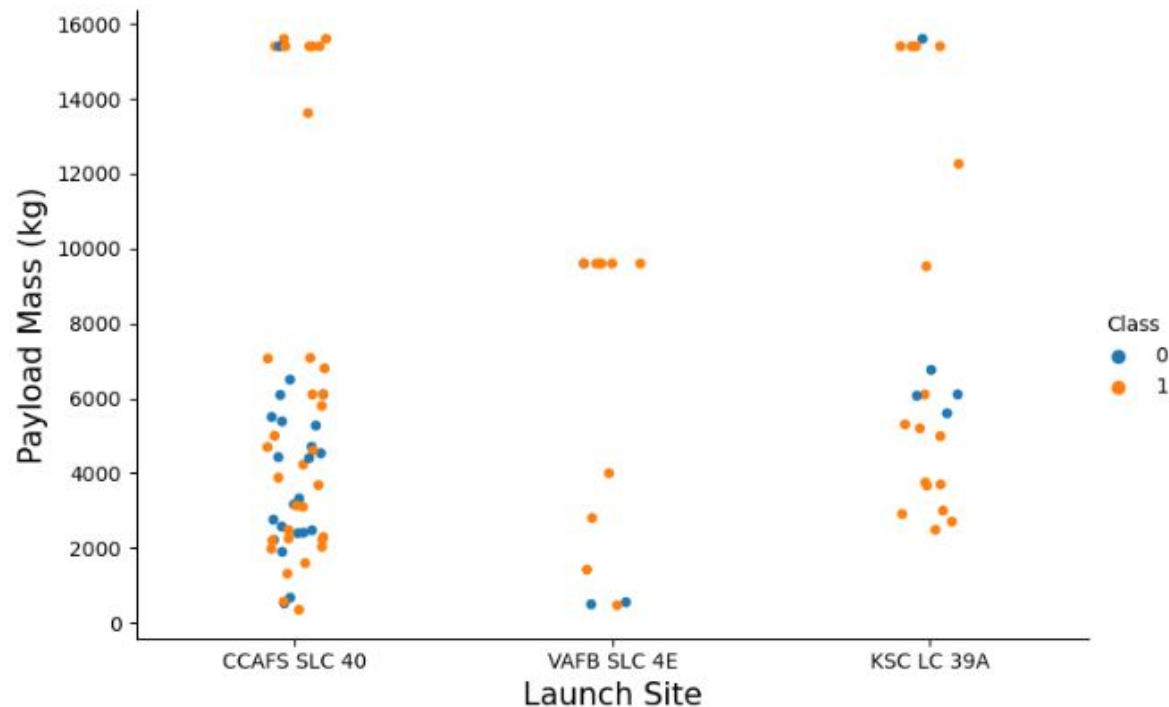
# Results: Exploratory Data Analysis

## Visualization and Feature Engineering

### TASK 2: Visualize the relationship between Payload and Launch Site

We also want to observe if there is any relationship between launch sites and their payload mass.

```
[7]: # Plot a scatter point chart with x axis to be Pay Load Mass (kg) and y axis to be the launch site, and hue to be the class value
sns.catplot(y="PayloadMass", x="LaunchSite", hue="Class", data=df, aspect = 2)
plt.xlabel("Launch Site", fontsize=15)
plt.ylabel("Payload Mass (kg)", fontsize=15)
plt.show()
```



# Results: Exploratory Data Analysis Visualization and Feature Engineering

## TASK 3: Visualize the relationship between success rate of each orbit type

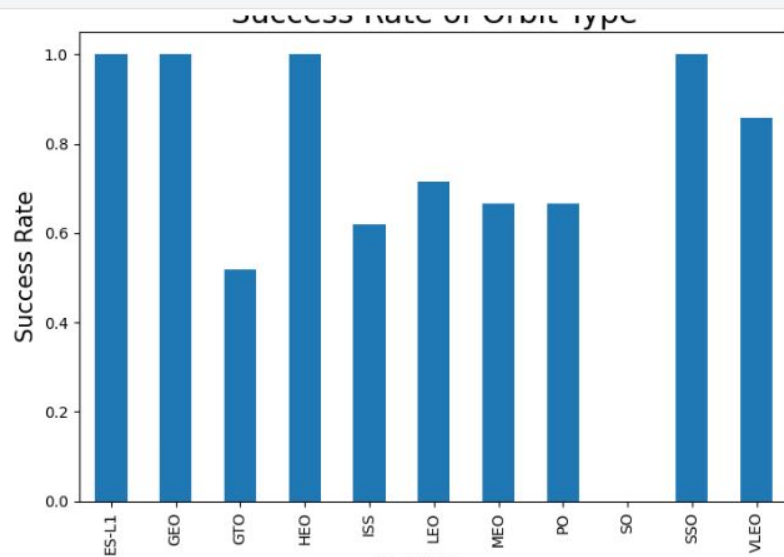
Next, we want to visually check if there are any relationship between success rate and orbit type.

Let's create a `bar chart` for the success rate of each orbit

```
[32]: # HINT use groupby method on Orbit column and get the mean of Class column

df.groupby('Orbit').mean()['Class'].plot(kind='bar')
plt.title('Success Rate of Orbit Type', fontsize=20)
plt.xlabel('Orbit Type', fontsize=15)
plt.ylabel('Success Rate', fontsize=15)

plt.show()
```





# Results: Exploratory Data Analysis

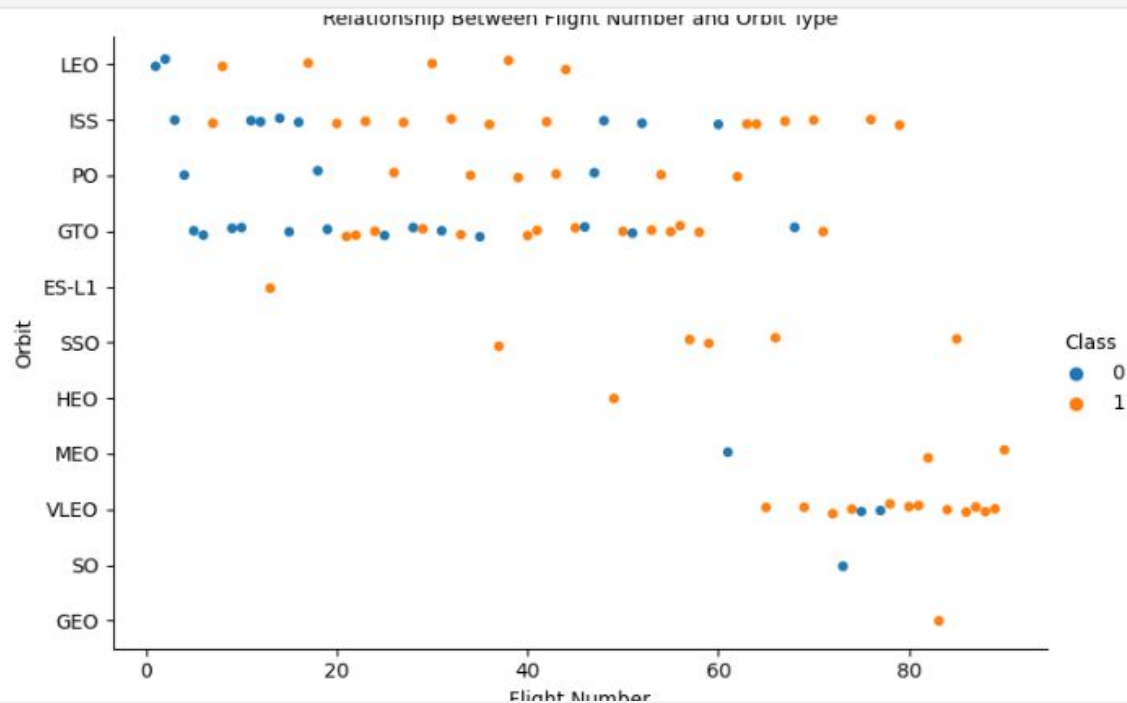
## Visualization and Feature Engineering

### TASK 4: Visualize the relationship between FlightNumber and Orbit type

For each orbit, we want to see if there is any relationship between FlightNumber and Orbit type.

```
[38]: # Plot a scatter point chart with x axis to be FlightNumber and y axis to be the Orbit, and hue to be the class value

sns.catplot(x='FlightNumber', y='Orbit', data=df, hue='Class', aspect = 1.5)
plt.title('Relationship Between Flight Number and Orbit Type', fontsize=10)
plt.xlabel('Flight Number', fontsize=10)
plt.show()
```



# Results: Exploratory Data Analysis


## Visualization and Feature Engineering

### TASK 5: Visualize the relationship between Payload and Orbit type

Similarly, we can plot the Payload vs. Orbit scatter point charts to reveal the relationship between Payload and Orbit type

```
10]: # Plot a scatter point chart with x axis to be Payload and y axis to be the Orbit, and hue to be the class value

sns.catplot(x='PayloadMass', y='Orbit', data=df, hue='Class', aspect = 5)
plt.title('Relationship between Payload Mass (Kg) and Orbit Type', fontsize=30)
plt.xlabel('Payload Mass (Kg)', fontsize=20)
plt.ylabel('Orbit Type', fontsize=20)
plt.show()
```



# Results: Exploratory Data Analysis Visualization and Feature Engineering

## TASK 6: Visualize the launch success yearly trend

You can plot a line chart with x axis to be `Year` and y axis to be average success rate, to get the average launch success trend.

The function will help you get the year from the date:

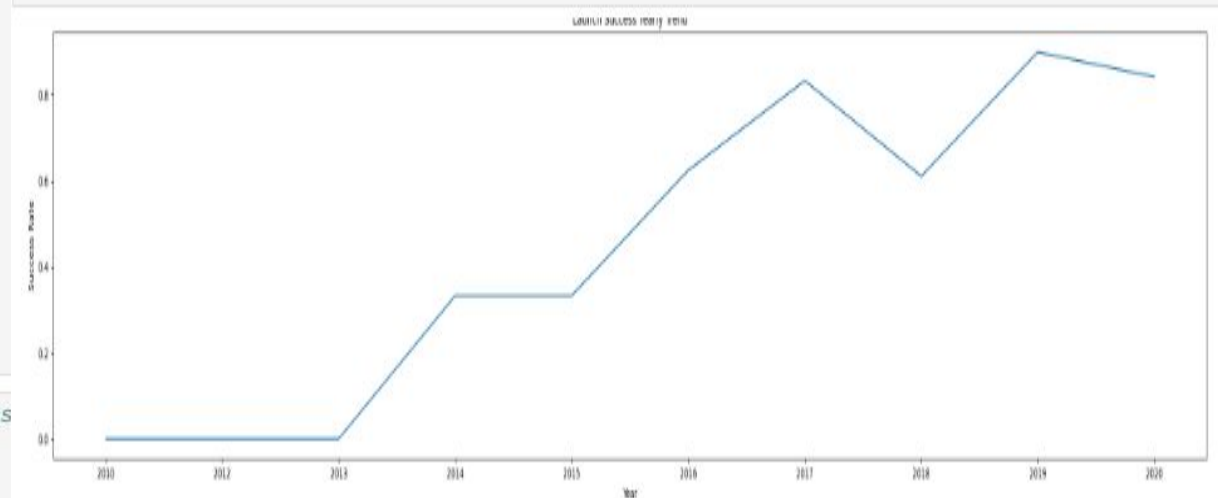
```
[11]: # A function to Extract years from the date
year=[]
def Extract_year(date):
    for i in df["Date"]:
        year.append(i.split("-")[0])
    return year

Extract_year(1)
df['Year']=year

avg_year=df.groupby(by='Year').mean()
avg_year.reset_index(inplace=True)

[12]: # Plot a line chart with x axis to be the extracted year and y axis to be the success

plt.plot(avg_year['Year'], avg_year['Class'])
plt.title('Launch Success Yearly Trend')
plt.xlabel('Year')
plt.ylabel('Success Rate')
plt.show()
```



you can observe that the success rate since 2013 kept increasing till 2020

# Results: Exploratory Data Analysis Visualization and Feature Engineering

## TASK 7: Create dummy variables to categorical columns

Use the function `get_dummies` and `features` dataframe to apply OneHotEncoder to the column `Orbits`, `LaunchSite`, `LandingPad`, and `Serial`. Assign the value to the variable `features_one_hot`, display the results using the method `head`. Your result dataframe must include all features including the encoded ones.

```
[44]: # HINT: Use get_dummies() function on the categorical columns
```

```
features_one_hot = pd.get_dummies(features, columns = ['Orbit', 'LaunchSite', 'LandingPad', 'Serial'])
features_one_hot.head(10)
```

```
[44]:
```

	FlightNumber	PayloadMass	Flights	GridFins	Reused	Legs	Block	ReusedCount	Orbit_ES- L1	Orbit_GEO	...	Serial_B1048	Serial_B1049	Serial_B1050	Serial_B:
0	1	6104.959412	1	False	False	False	1.0	0	0	0	...	0	0	0	
1	2	525.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	
2	3	677.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	
3	4	500.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	
4	5	3170.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	
5	6	3325.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	
6	7	2296.000000	1	False	False	True	1.0	0	0	0	...	0	0	0	
7	8	1316.000000	1	False	False	True	1.0	0	0	0	...	0	0	0	
8	9	4535.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	
9	10	4428.000000	1	False	False	False	1.0	0	0	0	...	0	0	0	

10 rows × 80 columns

# Results: Exploratory Data Analysis

## Visualization and Feature Engineering

TASK 8: Cast all numeric columns to `float64`

Now that our `features_one_hot` dataframe only contains numbers cast the entire dataframe to variable type `float64`

```
[5]: # HINT: use astype function
```

```
features_one_hot.astype('float64')
```

```
[5]:
```

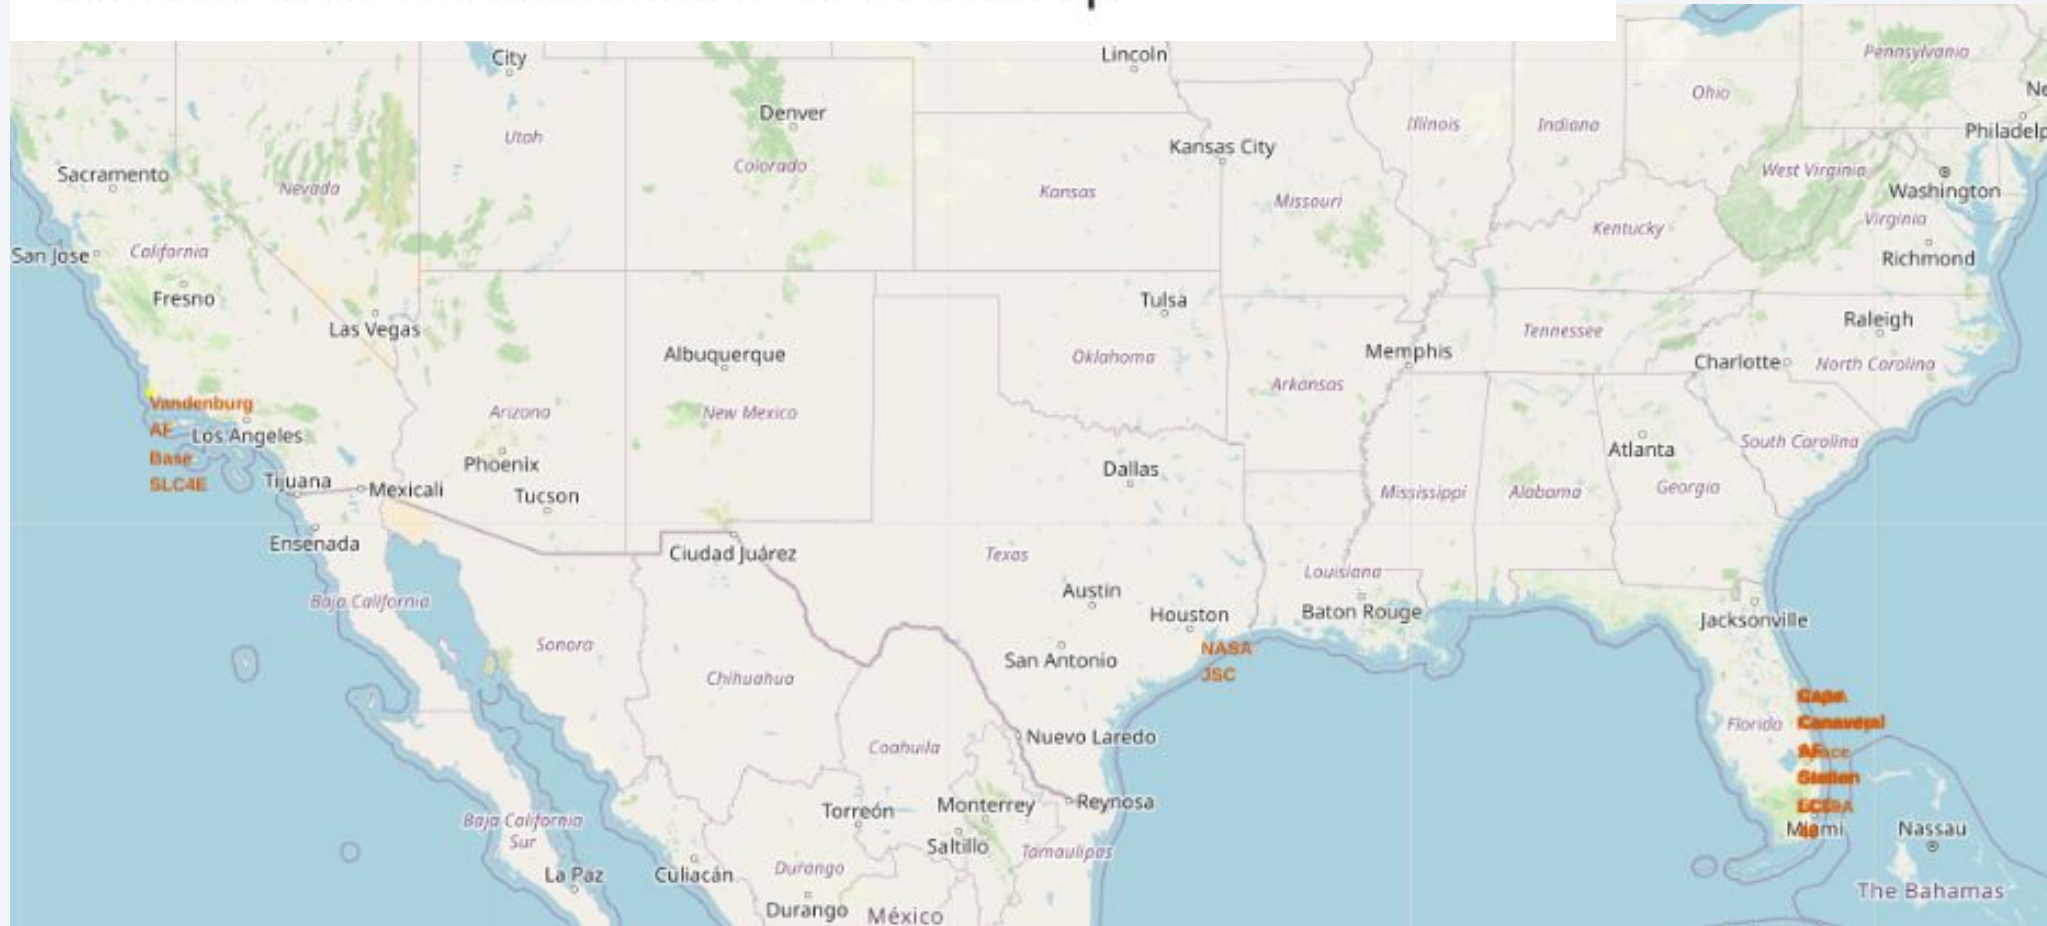
	FlightNumber	PayloadMass	Flights	GridFins	Reused	Legs	Block	ReusedCount	Orbit_ES-L1	Orbit_GEO	...	Serial_B1048	Serial_B1049	Serial_B1050	Serial_B1
0	1.0	6104.959412	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
1	2.0	525.000000	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
2	3.0	677.000000	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
3	4.0	500.000000	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
4	5.0	3170.000000	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
85	86.0	15400.000000	2.0	1.0	1.0	1.0	5.0	2.0	0.0	0.0	...	0.0	0.0	0.0	
86	87.0	15400.000000	3.0	1.0	1.0	1.0	5.0	2.0	0.0	0.0	...	0.0	0.0	0.0	
87	88.0	15400.000000	6.0	1.0	1.0	1.0	5.0	5.0	0.0	0.0	...	0.0	0.0	0.0	
88	89.0	15400.000000	3.0	1.0	1.0	1.0	5.0	2.0	0.0	0.0	...	0.0	0.0	0.0	
89	90.0	3681.000000	1.0	1.0	0.0	1.0	5.0	0.0	0.0	0.0	...	0.0	0.0	0.0	

90 rows × 80 columns



# Results: Interactive Maps with Folium

Task 1: Mark all launch sites on a map




Interactive map with launch sites and NASA JSC



# Results: Interactive Maps with Folium


Task 2: Mark the success/failed launches for each site on the map



Interactive map with KSC LC-40 Launches

# Results: Interactive Maps with Folium

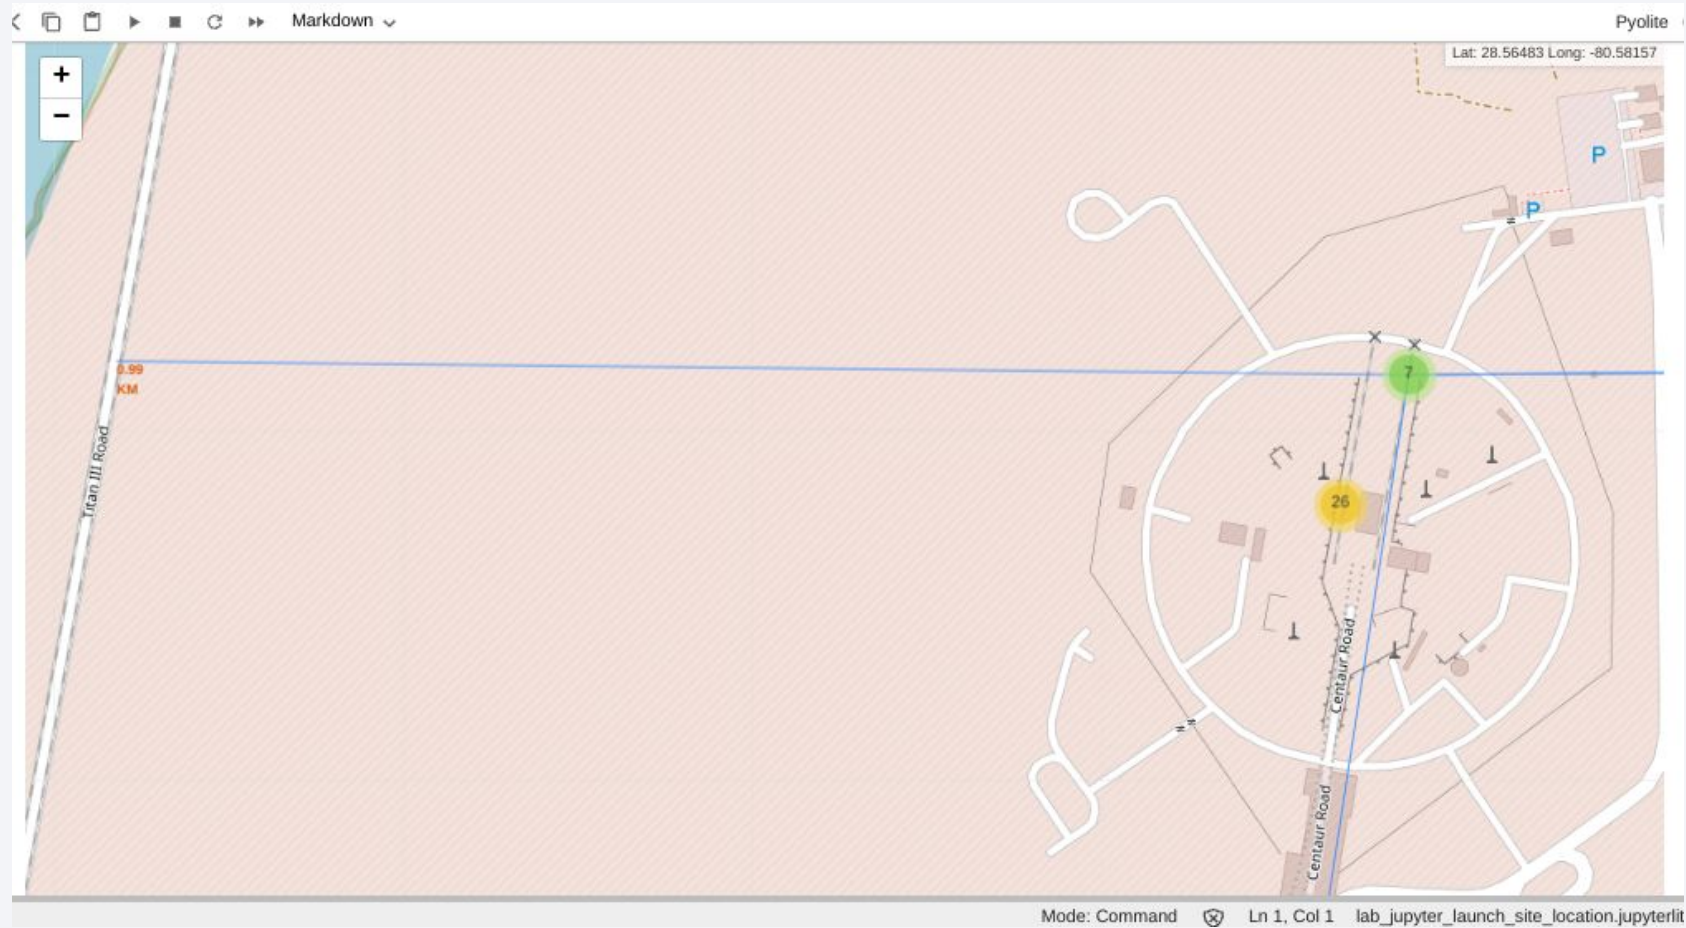
Task 2: Mark the success/failed launches for each site on the map



Interactive map with KSC SLC-40 Launches

# Results: Interactive Maps with Folium

TASK 3: Calculate the distances between a launch site to its proximities

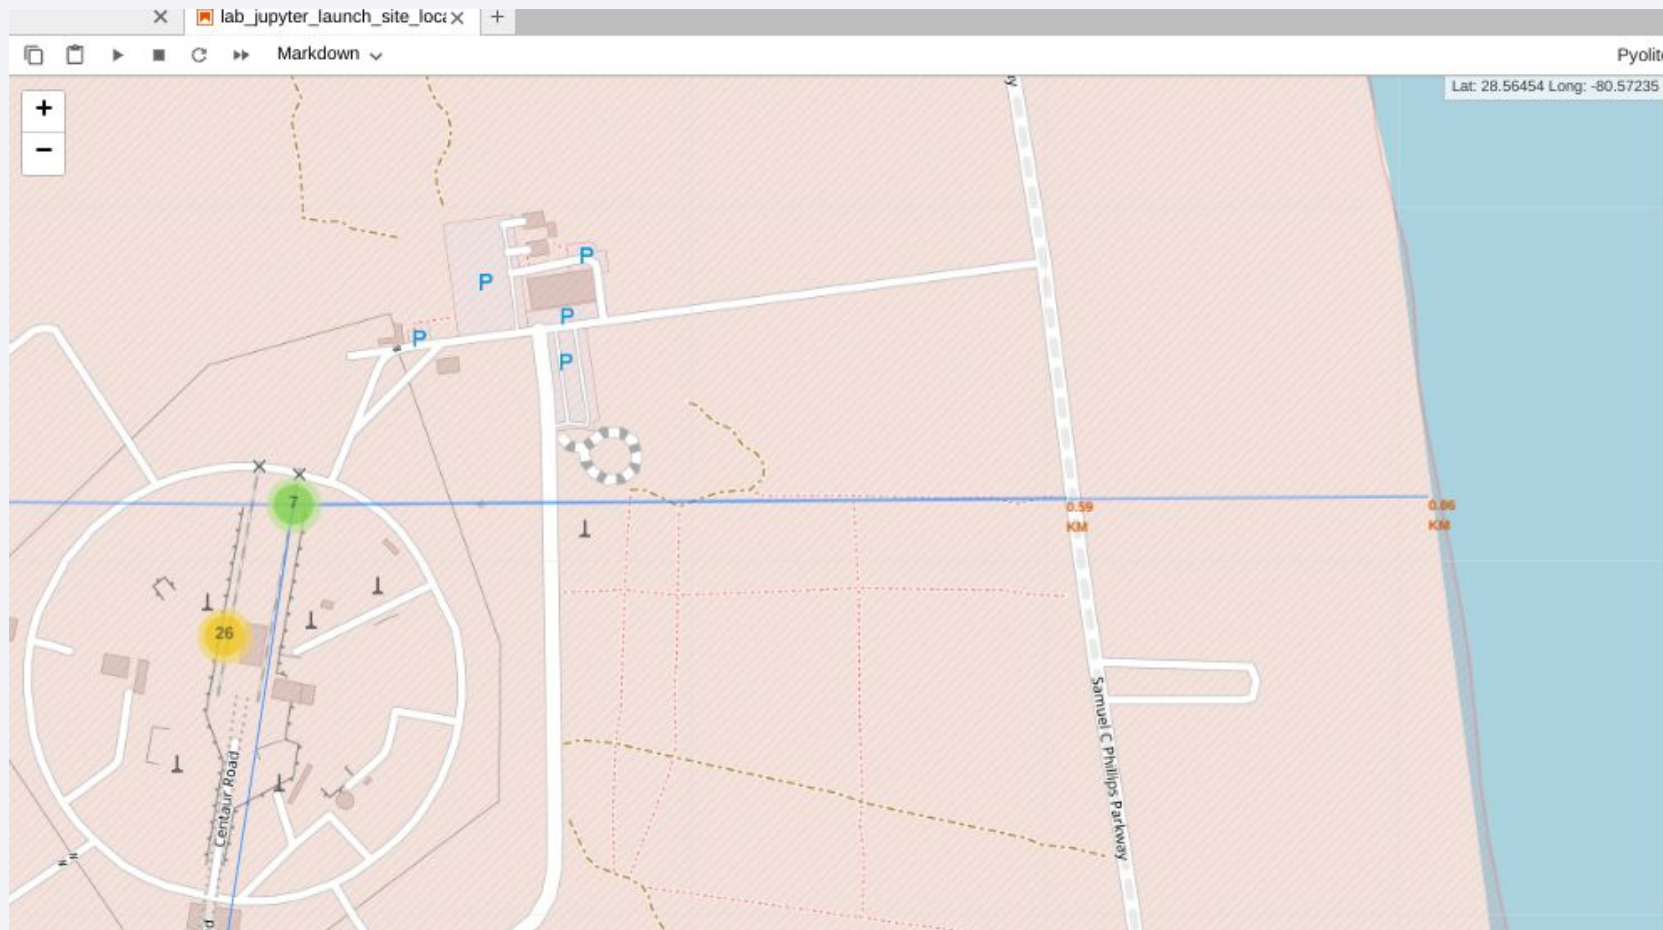


Interactive map with SLC-40 distance to closest railway (along Titan III Road)



# Results: Interactive Maps with Folium

## TASK 3: Calculate the distances between a launch site to its proximities




Interactive map with SLC-40 distance to closest road (Samuel C Phillips Parkway) and coastline

# Results: Interactive Maps with Folium

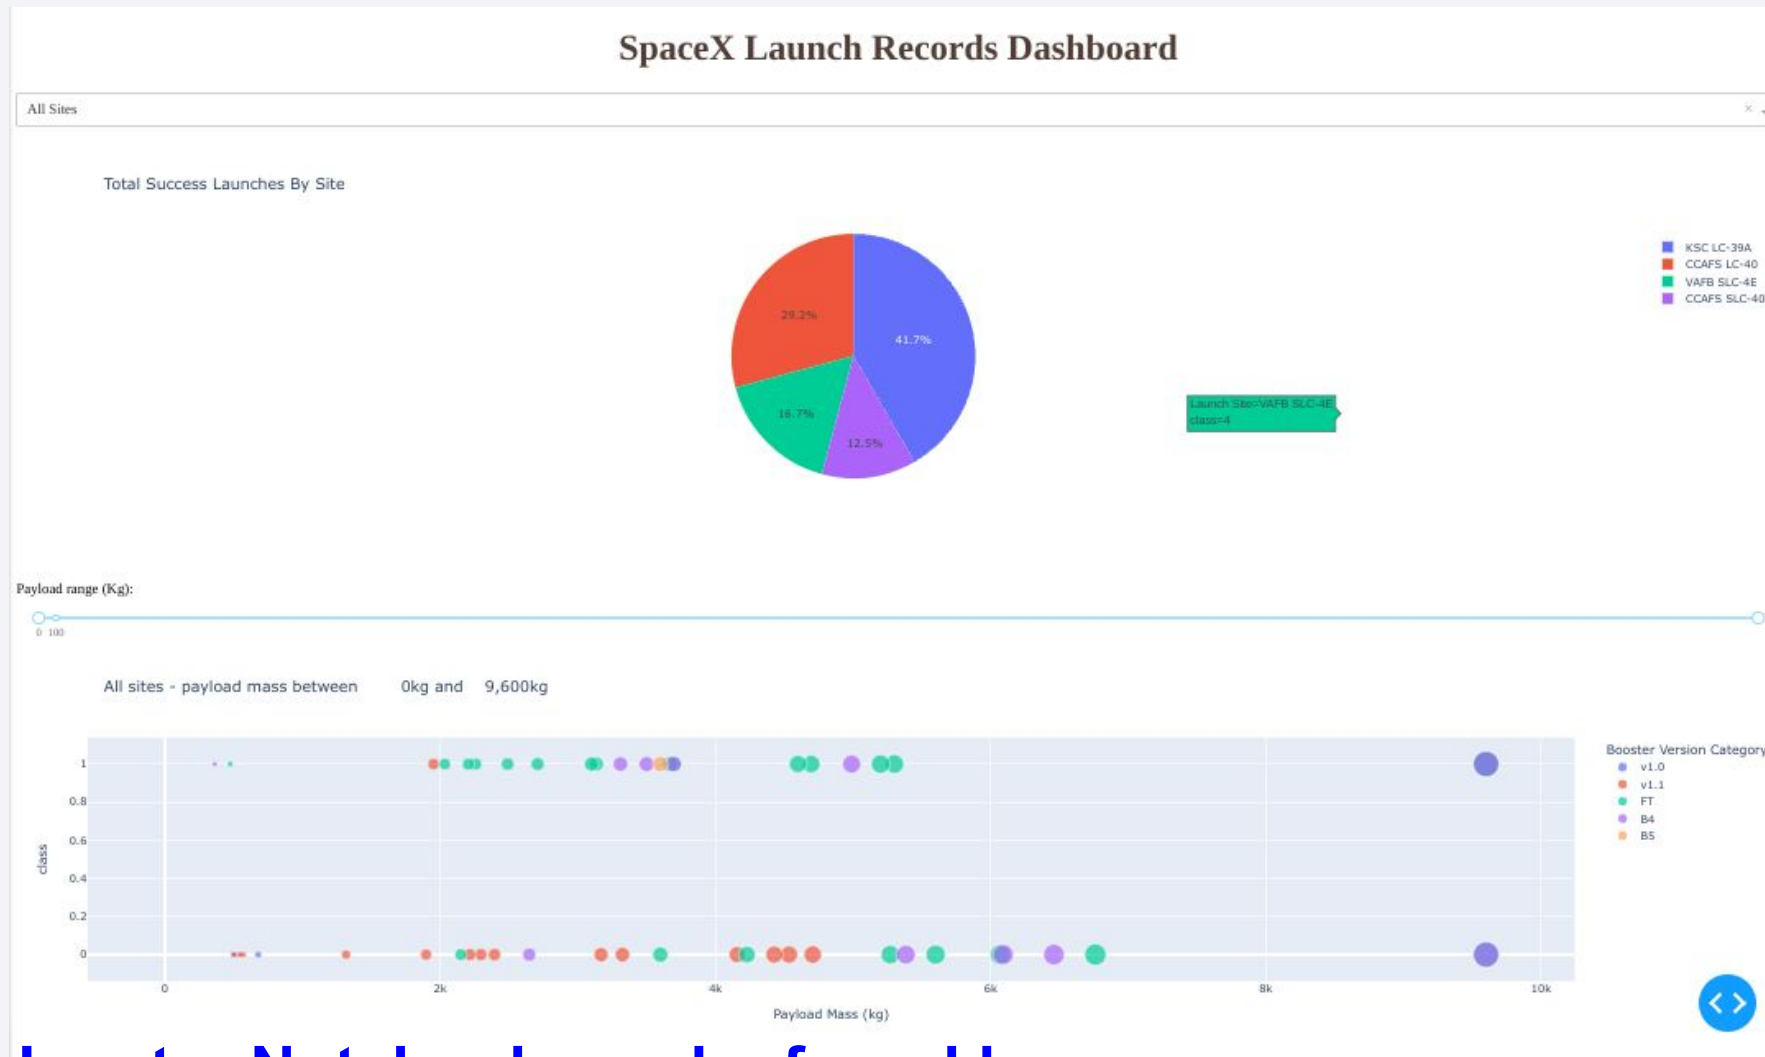
---

TASK 3: Calculate the distances between a launch site to its proximities



Interactive map with SLC-40 distance to closest city (Cape Canaveral, FL)

# Results: Dashboarding with Plotly Dash



Jupyter Notebook can be found here:



# Results: Predictive Analysis

## TASK 1

Create a NumPy array from the column `Class` in `data`, by applying the method `to_numpy()` then assign it to the variable `Y`, make sure the output is a Pandas series (only one bracket `df[name of column]`).

```
[8]: Y = data['Class'].to_numpy()
print(Y)

[0 0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 1 0 0 1 1 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1
 1 1 1 1 1 1 1 1 0 0 0 1 1 0 0 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 0 1 1 1 1 0 1
 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1]
```

## TASK 2

Standardize the data in `X` then reassign it to the variable `X` using the transform provided below.

```
[9]: # students get this
transform = preprocessing.StandardScaler()

[10]: X = transform.fit_transform(X)
print(X)

[[-1.71291154e+00 -1.94814463e-16 -6.53912840e-01 ... -8.35531692e-01
  1.93309133e+00 -1.93309133e+00]
 [-1.67441914e+00 -1.19523159e+00 -6.53912840e-01 ... -8.35531692e-01
  1.93309133e+00 -1.93309133e+00]
 [-1.63592675e+00 -1.16267307e+00 -6.53912840e-01 ... -8.35531692e-01
  1.93309133e+00 -1.93309133e+00]
 ...
 [ 1.63592675e+00  1.99100483e+00  3.49060516e+00 ...  1.19684269e+00
 -5.17306132e-01  5.17306132e-01]
 [ 1.67441914e+00  1.99100483e+00  1.00389436e+00 ...  1.19684269e+00
 -5.17306132e-01  5.17306132e-01]
 [ 1.71291154e+00 -5.19213966e-01 -6.53912840e-01 ... -8.35531692e-01
 -5.17306132e-01  5.17306132e-01]]
```

We split the data into training and testing data using the function `train_test_split`. The training data is divided into validation data, a second set used for training data; then the models are trained and hyperparameters are selected using the function `GridSearchCV`.

# Results: Predictive Analysis

---

## TASK 3

Use the function `train_test_split` to split the data `X` and `Y` into training and test data. Set the parameter `test_size` to 0.2 and `random_state` to 2. The training data and test data should be assigned to the following labels.

```
X_train, X_test, Y_train, Y_test
```

```
[11]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state = 2)
```

we can see we only have 18 test samples.

```
[12]: Y_test.shape
```

```
[12]: (18,)
```

# Results: Predictive Analysis

## TASK 4

Create a logistic regression object then create a GridSearchCV object `logreg_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
[13]: parameters = {'C': [0.01, 0.1, 1],  
                  'penalty': ['l2'],  
                  'solver': ['lbfgs']}  
  
[14]: lr_parameters = {"C": [0.01, 0.1, 1], 'penalty': ['l2'], 'solver': ['lbfgs']}# l1 lasso l2 ridge  
lr=LogisticRegression()  
logreg_cv = GridSearchCV(lr, lr_parameters, cv=10)  
logreg_cv.fit(X_train, Y_train)
```

```
[14]: GridSearchCV(cv=10, estimator=LogisticRegression(),  
                  param_grid={'C': [0.01, 0.1, 1], 'penalty': ['l2'],  
                              'solver': ['lbfgs']})
```

We output the `GridSearchCV` object for logistic regression. We display the best parameters using the data attribute `best_params_` and the accuracy on the validation data using the data attribute `best_score_`.

```
[15]: print("tuned hyperparameters :(best parameters) ",logreg_cv.best_params_)  
print("accuracy :",logreg_cv.best_score_)  
  
tuned hyperparameters :(best parameters) {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}  
accuracy : 0.8464285714285713
```

# Results: Predictive Analysis


## TASK 5

Calculate the accuracy on the test data using the method `score` :

```
[16]: print('Logistic Regression Test Accuracy Score :', logreg_cv.score(X_test, Y_test))
```

Logistic Regression Test Accuracy Score : 0.8333333333333334

```
[17]: yhat=logreg_cv.predict(X_test)
      plot_confusion_matrix(Y_test,yhat)
```



Examining the confusion matrix, we see that logistic regression can distinguish between the different classes. We see that the major problem is false positives.

# Results: Predictive Analysis

## TASK 6

Create a support vector machine object then create a `GridSearchCV` object `svm_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
[18]: svm_parameters = {'kernel':('linear', 'rbf', 'poly', 'rbf', 'sigmoid'),  
                        'C': np.logspace(-3, 3, 5),  
                        'gamma':np.logspace(-3, 3, 5)}  
svm = SVC()
```

```
[19]: svm_cv = GridSearchCV(svm, svm_parameters, cv=10)  
svm_cv.fit(X_train, Y_train)
```

```
[19]: GridSearchCV(cv=10, estimator=SVC(),  
                  param_grid={'C': array([1.00000000e-03, 3.16227766e-02, 1.00000000e+00, 3.16227766e+01,  
1.00000000e+03]),  
                              'gamma': array([1.00000000e-03, 3.16227766e-02, 1.00000000e+00, 3.16227766e+01,  
1.00000000e+03]),  
                              'kernel': ('linear', 'rbf', 'poly', 'rbf', 'sigmoid')})
```

```
[20]: print("tuned hpyerparameters :(best parameters) ",svm_cv.best_params_)  
print("accuracy :",svm_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'C': 1.0, 'gamma': 0.03162277660168379, 'kernel': 'sigmoid'}  
accuracy : 0.8482142857142856
```

# Results: Predictive Analysis


## TASK 7

Calculate the accuracy on the test data using the method `score` :

```
[21]: print('SVM Test Accuracy Score :', svm_cv.score(X_test, Y_test))
```

```
SVM Test Accuracy Score : 0.8333333333333334
```

```
[22]: yhat=svm_cv.predict(X_test)  
plot_confusion_matrix(Y_test,yhat)
```



# Results: Predictive Analysis

## TASK 8

Create a decision tree classifier object then create a `GridSearchCV` object `tree_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
[35]: tree_parameters = {'criterion': ['gini', 'entropy'],
                        'splitter': ['best', 'random'],
                        'max_depth': [2*n for n in range(1,10)],
                        'max_features': ['auto', 'sqrt'],
                        'min_samples_leaf': [1, 2, 4],
                        'min_samples_split': [2, 5, 10]}

tree = DecisionTreeClassifier()
```

```
[36]: tree_cv = GridSearchCV(tree, tree_parameters, cv=10)
tree_cv.fit(X_train, Y_train)
```

```
[36]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(),
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18],
                              'max_features': ['auto', 'sqrt'],
                              'min_samples_leaf': [1, 2, 4],
                              'min_samples_split': [2, 5, 10],
                              'splitter': ['best', 'random']})
```

```
[37]: print("tuned hyperparameters :(best parameters) ",tree_cv.best_params_)
print("accuracy :",tree_cv.best_score_)
```

```
tuned hyperparameters :(best parameters) {'criterion': 'entropy', 'max_depth': 2, 'max_features': 'auto', 'min_samples_leaf': 2, 'min_samples_split': 2, 'splitter': 'random'}
accuracy : 0.875
```



# Results: Predictive Analysis

## TASK 9


Calculate the accuracy of tree\_cv on the test data using the method `score` :

```
[38]: print('Decision Tree Test Accuracy Score :', tree_cv.score(X_test, Y_test))
```

Decision Tree Test Accuracy Score : 0.8333333333333334

We can plot the confusion matrix

```
[39]: yhat = svm_cv.predict(X_test)
      plot_confusion_matrix(Y_test, yhat)
```



# Results: Predictive Analysis

## TASK 10

Create a k nearest neighbors object then create a `GridSearchCV` object `knn_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
[40]: KNN_parameters = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],  
                        'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],  
                        'p': [1, 2]}
```

```
KNN = KNeighborsClassifier()
```

```
[41]: knn_cv = GridSearchCV(KNN, KNN_parameters, cv=10)  
      knn_cv.fit(X_train, Y_train)
```

```
[41]: GridSearchCV(cv=10, estimator=KNeighborsClassifier(),  
                  param_grid={'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],  
                              'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],  
                              'p': [1, 2]})
```

```
[42]: print("tuned hpyerparameters :(best parameters) ",knn_cv.best_params_)  
      print("accuracy :",knn_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'algorithm': 'auto', 'n_neighbors': 10, 'p': 1}  
accuracy : 0.8482142857142858
```

# Results: Predictive Analysis

## TASK 11


Calculate the accuracy of knn\_cv on the test data using the method `score` :

```
: print('KNN Test Accuracy Score :', knn_cv.score(X_test, Y_test))
```

KNN Test Accuracy Score : 0.8333333333333334

We can plot the confusion matrix

```
[44]: yhat = knn_cv.predict(X_test)  
      plot_confusion_matrix(Y_test, yhat)
```



# Results: Predictive Analysis

---

## TASK 12

Find the method performs best:

```
[45]: #all of the tests revealed all had the same accuracy score so all perform the same. Due to the small data set for each test.  
print('Logistic Regression Test Accuracy Score :', logreg_cv.score(X_test, Y_test))  
print('SVM Test Accuracy Score :', svm_cv.score(X_test, Y_test))  
print('Decision Tree Test Accuracy Score :', tree_cv.score(X_test, Y_test))  
print('KNN Test Accuracy Score :', knn_cv.score(X_test, Y_test))
```

```
Logistic Regression Test Accuracy Score : 0.8333333333333334  
SVM Test Accuracy Score : 0.8333333333333334  
Decision Tree Test Accuracy Score : 0.8333333333333334  
KNN Test Accuracy Score : 0.8333333333333334
```

# Conclusions

---

- Data was obtained and converted into useful forms
- Exploratory data analysis (EDA) carried out using SQL queries and visualizations in Seabor/Matplotlib
- Interactive visual analytics via Folium
- Dashboarding using Plotly Dash
- Predictive analysis was demonstrated for various classification models

# Appendix

---



Thank you!

